

all_auto

January 13, 2025

```
[1]: # use a library to do everything automatically, so that we can compare our
      ↪ results to

import numpy as np
import pandas as pd
import plotly.express as px
import plotly.graph_objects as go
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import statsmodels.tsa.stattools as st
import statsmodels.api as sm
import numpy as np
import pmdarima as pm

# Load the CSV file
data = pd.read_csv('../data/EUR_USD_monthly.csv')

# Extract the time and price series
time = pd.to_datetime(data.iloc[:, 0]) # Assuming first column is 'time'
prices = data.iloc[:, 1].replace(',', '', regex=True).astype(float) # Remove
      ↪ commas and convert to float

# Create a DataFrame with the time and prices
time_series_df = pd.DataFrame({'time': time, 'points': prices})
time_series_df.set_index('time', inplace=True)

def plot_series(series):
    # Plot the time series using Plotly
    series.plot()

    plot_acf(series)
    plot_pacf(series)
plot_series(time_series_df)
```

```

# Automatically fit the best ARIMA/SARIMA model
model = pm.auto_arima(
    time_series_df,
    seasonal=False,      # Enable seasonal differencing if necessary
    stepwise=False,      # Use stepwise algorithm for faster computation
    trace=True,          # Show the fitting process details
    #error_action='ignore', # Ignore if the model doesn't fit properly
    #suppress_warnings=True # Suppress warnings
    max_p=10,            # Set maximum p value to control AR component
    max_q=10,            # Set maximum q value to control MA component
    max_d=1,             # Set maximum differencing order
    max_P=1,             # Set maximum seasonal AR order
    max_Q=1,             # Set maximum seasonal MA order
    max_D=1,
    max_order = 10000
)

# Summary of the best model found
print(model.summary())

# Predict future values (e.g., next 12 periods)
forecast = model.predict(n_periods=12)
print(forecast)

```

```

ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-780.323, Time=0.05 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-778.566, Time=0.09 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-776.664, Time=0.12 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=-777.610, Time=0.20 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=-775.616, Time=0.37 sec
ARIMA(0,1,5)(0,0,0)[0] intercept : AIC=-773.677, Time=0.33 sec
ARIMA(0,1,6)(0,0,0)[0] intercept : AIC=-774.825, Time=0.49 sec
ARIMA(0,1,7)(0,0,0)[0] intercept : AIC=-786.213, Time=0.61 sec
ARIMA(0,1,8)(0,0,0)[0] intercept : AIC=-788.144, Time=0.86 sec
ARIMA(0,1,9)(0,0,0)[0] intercept : AIC=-788.023, Time=1.14 sec
ARIMA(0,1,10)(0,0,0)[0] intercept : AIC=-786.866, Time=1.38 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-778.554, Time=0.08 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-776.576, Time=0.20 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-775.663, Time=0.26 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=-772.678, Time=0.38 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=-773.612, Time=0.82 sec
ARIMA(1,1,5)(0,0,0)[0] intercept : AIC=-771.732, Time=0.63 sec
ARIMA(1,1,6)(0,0,0)[0] intercept : AIC=-772.217, Time=0.75 sec
ARIMA(1,1,7)(0,0,0)[0] intercept : AIC=-789.623, Time=1.25 sec
ARIMA(1,1,8)(0,0,0)[0] intercept : AIC=-787.402, Time=1.28 sec
ARIMA(1,1,9)(0,0,0)[0] intercept : AIC=-785.811, Time=0.91 sec
ARIMA(1,1,10)(0,0,0)[0] intercept : AIC=-784.508, Time=1.16 sec

```

```

ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=-776.734, Time=0.14 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-776.057, Time=0.39 sec
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-779.516, Time=0.38 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=-777.408, Time=0.41 sec
ARIMA(2,1,4)(0,0,0)[0] intercept : AIC=-776.261, Time=0.61 sec
ARIMA(2,1,5)(0,0,0)[0] intercept : AIC=-773.019, Time=0.74 sec
ARIMA(2,1,6)(0,0,0)[0] intercept : AIC=-771.591, Time=0.72 sec
ARIMA(2,1,7)(0,0,0)[0] intercept : AIC=-788.194, Time=0.89 sec
ARIMA(2,1,8)(0,0,0)[0] intercept : AIC=-786.177, Time=3.69 sec
ARIMA(2,1,9)(0,0,0)[0] intercept : AIC=-783.308, Time=1.20 sec
ARIMA(2,1,10)(0,0,0)[0] intercept : AIC=-783.433, Time=1.09 sec
ARIMA(3,1,0)(0,0,0)[0] intercept : AIC=-778.059, Time=0.15 sec
ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=-776.526, Time=0.38 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=-776.032, Time=0.60 sec
ARIMA(3,1,3)(0,0,0)[0] intercept : AIC=-775.393, Time=0.78 sec
ARIMA(3,1,4)(0,0,0)[0] intercept : AIC=-777.290, Time=0.68 sec
ARIMA(3,1,5)(0,0,0)[0] intercept : AIC=-778.172, Time=0.75 sec
ARIMA(3,1,6)(0,0,0)[0] intercept : AIC=-774.227, Time=0.65 sec
ARIMA(3,1,7)(0,0,0)[0] intercept : AIC=-785.189, Time=0.76 sec
ARIMA(3,1,8)(0,0,0)[0] intercept : AIC=-784.013, Time=1.03 sec
ARIMA(3,1,9)(0,0,0)[0] intercept : AIC=-780.698, Time=1.02 sec
ARIMA(3,1,10)(0,0,0)[0] intercept : AIC=-782.129, Time=1.07 sec
ARIMA(4,1,0)(0,0,0)[0] intercept : AIC=-776.713, Time=0.15 sec
ARIMA(4,1,1)(0,0,0)[0] intercept : AIC=-774.712, Time=0.20 sec
ARIMA(4,1,2)(0,0,0)[0] intercept : AIC=-772.716, Time=0.48 sec
ARIMA(4,1,3)(0,0,0)[0] intercept : AIC=-773.050, Time=0.56 sec
ARIMA(4,1,4)(0,0,0)[0] intercept : AIC=-775.124, Time=0.59 sec
ARIMA(4,1,5)(0,0,0)[0] intercept : AIC=inf, Time=0.72 sec
ARIMA(4,1,6)(0,0,0)[0] intercept : AIC=-772.004, Time=0.86 sec
ARIMA(4,1,7)(0,0,0)[0] intercept : AIC=-788.061, Time=1.14 sec
ARIMA(4,1,8)(0,0,0)[0] intercept : AIC=-786.468, Time=1.16 sec
ARIMA(4,1,9)(0,0,0)[0] intercept : AIC=-783.727, Time=1.09 sec
ARIMA(4,1,10)(0,0,0)[0] intercept : AIC=-779.023, Time=1.21 sec
ARIMA(5,1,0)(0,0,0)[0] intercept : AIC=-774.732, Time=0.19 sec
ARIMA(5,1,1)(0,0,0)[0] intercept : AIC=-772.734, Time=0.23 sec
ARIMA(5,1,2)(0,0,0)[0] intercept : AIC=-765.585, Time=0.52 sec
ARIMA(5,1,3)(0,0,0)[0] intercept : AIC=-770.939, Time=0.48 sec
ARIMA(5,1,4)(0,0,0)[0] intercept : AIC=-773.459, Time=0.59 sec
ARIMA(5,1,5)(0,0,0)[0] intercept : AIC=-772.102, Time=0.77 sec
ARIMA(5,1,6)(0,0,0)[0] intercept : AIC=-772.568, Time=0.81 sec
ARIMA(5,1,7)(0,0,0)[0] intercept : AIC=-784.035, Time=0.82 sec
ARIMA(5,1,8)(0,0,0)[0] intercept : AIC=-783.259, Time=0.95 sec
ARIMA(5,1,9)(0,0,0)[0] intercept : AIC=-778.983, Time=1.05 sec
ARIMA(5,1,10)(0,0,0)[0] intercept : AIC=-776.574, Time=1.14 sec
ARIMA(6,1,0)(0,0,0)[0] intercept : AIC=-773.346, Time=0.18 sec
ARIMA(6,1,1)(0,0,0)[0] intercept : AIC=-775.683, Time=0.55 sec
ARIMA(6,1,2)(0,0,0)[0] intercept : AIC=-771.604, Time=0.59 sec
ARIMA(6,1,3)(0,0,0)[0] intercept : AIC=-770.539, Time=1.09 sec

```

ARIMA(6,1,4)(0,0,0)[0] intercept : AIC=-777.565, Time=1.20 sec
 ARIMA(6,1,5)(0,0,0)[0] intercept : AIC=-763.117, Time=0.79 sec
 ARIMA(6,1,6)(0,0,0)[0] intercept : AIC=-770.514, Time=0.93 sec
 ARIMA(6,1,7)(0,0,0)[0] intercept : AIC=-781.841, Time=0.94 sec
 ARIMA(6,1,8)(0,0,0)[0] intercept : AIC=-780.520, Time=1.12 sec
 ARIMA(6,1,9)(0,0,0)[0] intercept : AIC=-780.710, Time=1.05 sec
 ARIMA(6,1,10)(0,0,0)[0] intercept : AIC=-776.810, Time=1.18 sec
 ARIMA(7,1,0)(0,0,0)[0] intercept : AIC=-780.185, Time=0.57 sec
 ARIMA(7,1,1)(0,0,0)[0] intercept : AIC=-784.702, Time=0.56 sec
 ARIMA(7,1,2)(0,0,0)[0] intercept : AIC=-785.095, Time=0.62 sec
 ARIMA(7,1,3)(0,0,0)[0] intercept : AIC=-778.585, Time=0.71 sec
 ARIMA(7,1,4)(0,0,0)[0] intercept : AIC=-778.842, Time=0.70 sec
 ARIMA(7,1,5)(0,0,0)[0] intercept : AIC=-780.412, Time=0.76 sec
 ARIMA(7,1,6)(0,0,0)[0] intercept : AIC=-782.625, Time=0.87 sec
 ARIMA(7,1,7)(0,0,0)[0] intercept : AIC=-782.239, Time=0.93 sec
 ARIMA(7,1,8)(0,0,0)[0] intercept : AIC=-778.164, Time=1.09 sec
 ARIMA(7,1,9)(0,0,0)[0] intercept : AIC=-778.036, Time=1.10 sec
 ARIMA(7,1,10)(0,0,0)[0] intercept : AIC=-772.525, Time=1.19 sec
 ARIMA(8,1,0)(0,0,0)[0] intercept : AIC=-780.928, Time=0.37 sec
 ARIMA(8,1,1)(0,0,0)[0] intercept : AIC=-783.543, Time=0.70 sec
 ARIMA(8,1,2)(0,0,0)[0] intercept : AIC=-784.090, Time=0.87 sec
 ARIMA(8,1,3)(0,0,0)[0] intercept : AIC=-778.282, Time=0.88 sec
 ARIMA(8,1,4)(0,0,0)[0] intercept : AIC=-778.121, Time=0.92 sec
 ARIMA(8,1,5)(0,0,0)[0] intercept : AIC=-780.449, Time=1.01 sec
 ARIMA(8,1,6)(0,0,0)[0] intercept : AIC=-781.395, Time=1.02 sec
 ARIMA(8,1,7)(0,0,0)[0] intercept : AIC=-779.820, Time=0.99 sec
 ARIMA(8,1,8)(0,0,0)[0] intercept : AIC=-779.167, Time=1.15 sec
 ARIMA(8,1,9)(0,0,0)[0] intercept : AIC=-777.241, Time=1.24 sec
 ARIMA(8,1,10)(0,0,0)[0] intercept : AIC=-776.783, Time=1.42 sec
 ARIMA(9,1,0)(0,0,0)[0] intercept : AIC=-784.241, Time=0.50 sec
 ARIMA(9,1,1)(0,0,0)[0] intercept : AIC=-783.348, Time=0.75 sec
 ARIMA(9,1,2)(0,0,0)[0] intercept : AIC=-781.308, Time=0.80 sec
 ARIMA(9,1,3)(0,0,0)[0] intercept : AIC=-781.580, Time=0.97 sec
 ARIMA(9,1,4)(0,0,0)[0] intercept : AIC=-781.300, Time=0.92 sec
 ARIMA(9,1,5)(0,0,0)[0] intercept : AIC=-783.088, Time=1.02 sec
 ARIMA(9,1,6)(0,0,0)[0] intercept : AIC=-780.519, Time=1.17 sec
 ARIMA(9,1,7)(0,0,0)[0] intercept : AIC=-776.893, Time=1.19 sec
 ARIMA(9,1,8)(0,0,0)[0] intercept : AIC=-776.257, Time=1.29 sec
 ARIMA(9,1,9)(0,0,0)[0] intercept : AIC=-774.141, Time=1.30 sec
 ARIMA(9,1,10)(0,0,0)[0] intercept : AIC=-774.679, Time=1.49 sec
 ARIMA(10,1,0)(0,0,0)[0] intercept : AIC=-783.369, Time=0.54 sec
 ARIMA(10,1,1)(0,0,0)[0] intercept : AIC=-781.589, Time=0.89 sec
 ARIMA(10,1,2)(0,0,0)[0] intercept : AIC=-780.122, Time=0.99 sec
 ARIMA(10,1,3)(0,0,0)[0] intercept : AIC=-781.219, Time=1.06 sec
 ARIMA(10,1,4)(0,0,0)[0] intercept : AIC=-779.766, Time=1.24 sec
 ARIMA(10,1,5)(0,0,0)[0] intercept : AIC=-777.893, Time=1.56 sec
 ARIMA(10,1,6)(0,0,0)[0] intercept : AIC=-776.533, Time=1.64 sec
 ARIMA(10,1,7)(0,0,0)[0] intercept : AIC=-776.063, Time=1.23 sec

ARIMA(10,1,8)(0,0,0)[0] intercept : AIC=-774.097, Time=1.40 sec
 ARIMA(10,1,9)(0,0,0)[0] intercept : AIC=-772.588, Time=1.56 sec
 ARIMA(10,1,10)(0,0,0)[0] intercept : AIC=-772.297, Time=2.22 sec

Best model: ARIMA(1,1,7)(0,0,0)[0] intercept
 Total fit time: 102.259 seconds

SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:          204
Model:                 SARIMAX(1, 1, 7)      Log Likelihood          404.811
Date:                 Mon, 13 Jan 2025      AIC                  -789.623
Time:                 18:44:13      BIC                  -756.491
Sample:              0      HQIC                  -776.219
                        - 204
Covariance Type:      opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	0.0012	0.000	2.504	0.012	0.000	0.002
ar.L1	0.4156	0.173	2.399	0.016	0.076	0.755
ma.L1	-0.4756	0.158	-3.011	0.003	-0.785	-0.166
ma.L2	-0.0424	0.063	-0.677	0.498	-0.165	0.080
ma.L3	0.1395	0.061	2.269	0.023	0.019	0.260
ma.L4	-0.1904	0.061	-3.124	0.002	-0.310	-0.071
ma.L5	-0.0169	0.059	-0.285	0.775	-0.133	0.099
ma.L6	0.0529	0.077	0.687	0.492	-0.098	0.204
ma.L7	-0.3274	0.068	-4.804	0.000	-0.461	-0.194
sigma2	0.0011	8.76e-05	12.293	0.000	0.001	0.001

```
=====
Ljung-Box (L1) (Q):          0.00      Jarque-Bera (JB):
22.66
Prob(Q):                    0.96      Prob(JB):
0.00
Heteroskedasticity (H):      4.02      Skew:
0.00
Prob(H) (two-sided):        0.00      Kurtosis:
4.64
=====
```

Warnings:


[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```
204    1.443172
205    1.438601
206    1.442985
207    1.438519
```


```
208    1.450690
209    1.456158
210    1.459125
211    1.461514
212    1.463662
213    1.465711
214    1.467718
215    1.469708
dtype: float64
```

```
/home/daniels/.conda/envs/adlr/lib/python3.10/site-
packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning: No supported index
is available. Prediction results will be given with an integer index beginning
at `start`.
```

```
    return get_prediction_index(
/home/daniels/.conda/envs/adlr/lib/python3.10/site-
packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning: No supported
index is available. In the next version, calling this method in a model without
a supported index will result in an exception.
    return get_prediction_index(
```







[]:

Test stationarity, Augmented Dickey-Fuller unit root test.

```
[2]: #Check stationarity
# alternative hypothesis is stationarity
st.adfuller(time_series_df['points'])
```

```
[2]: (-0.8833017572640469,
0.7934216416235764,
9,
194,
{'1%': -3.4645146202692527,
'5%': -2.8765564361715534,
'10%': -2.5747745328940375},
-717.875322423628)
```

p value > 0.05 => we can not conclude stationarity

[]:

apply box cox transformation, and use scipy to estimate optimal box cox parameter lambda (todo explain what method scipy uses)


```
[3]: from scipy import stats


# Apply Box-Cox transformation and estimate lambda
transformed_prices, lambda_est = stats.boxcox(time_series_df["points"])

# Store the transformed values back into the DataFrame
time_series_df['boxcox_points'] = transformed_prices

time_series_df['boxcox_points'].plot()

print("Optimal lambda: ", lambda_est)
```

Optimal lambda: -2.3018585777936003



Differentiate and see when stationarity tests become positive

```
[4]: alpha = 0.05

# Assuming 'time_series_df' is your DataFrame and 'Prices' is the column
```

```

current_series = time_series_df['points'] # todo possible bug, try with_: pd.
↳ DataFrame(np.random.random(100))#np.exp(time_series_df['boxcox_points'])
differencing_count = 0

# Function to apply ADF test and get p-value
def get_adf_p_value(series):
    adf_result = st.adfuller(series.dropna())
    return adf_result[1] # Return p-value from ADF test

# Keep iterating until the series is stationary or we reach the max differencing
while get_adf_p_value(current_series) >= alpha:
    differencing_count += 1
    # Apply first-order differencing
    current_series = current_series.diff()

    # Add the differenced series to the DataFrame
    time_series_df[f'diff_prices_{differencing_count}'] = current_series
    print("defferentiation")
    # Print progress with the p-value after differencing
    print(f"After {differencing_count} differencing(s), the ADF p-value is:↳
↳ {get_adf_p_value(current_series)}")

# Check if series is stationary after differencing
if get_adf_p_value(current_series) < alpha:
    print(f"Series is stationary after {differencing_count} differencings.")
    print("")
else:
    print("Max differencings reached without achieving stationarity.")

```


defferentiation


After 1 differencing(s), the ADF p-value is: 1.2426670145880777e-08


Series is stationary after 1 differencings.

Plot the new time series and acf and pcf

```
[5]: plot_series(time_series_df[f'diff_prices_{differencing_count}'].dropna())
```







Print significant auto and partial autocorrelations

```
[6]: print(st.acf(time_series_df[f'diff_prices_{differencing_count}'].dropna(),
    ↪qstat=True))
```

```
(array([ 1.          , -0.03365732, -0.02826342,  0.12864677, -0.06457575,
        0.00608376,  0.07312028, -0.22007495, -0.09050825, -0.10975887,
        0.04149674, -0.08338499, -0.12708712, -0.05490398, -0.03557711,
        -0.02422668,  0.11440957,  0.00193089,  0.05805922,  0.12820226,
        -0.05282868,  0.06202194, -0.01250253, -0.0767693 ]), array([ 0.23337674,
        0.39876442,  3.84240378,  4.71444248,  4.72222158,
        5.85165179, 16.1350117 , 17.88321363, 20.46742119, 20.83871736,
        22.34575616, 25.86475895, 26.52500207, 26.8036972 , 26.9336183 ,
        29.84656821, 29.84740237, 30.60566648, 34.3229306 , 34.95758643,
        35.83715387, 35.87309297, 37.23564319]), array([0.62903125, 0.81923671,
        0.27899224, 0.31787164, 0.45071202,
        0.44001183, 0.0239111 , 0.0221193 , 0.01523701, 0.0222467 ,
        0.02182736, 0.01121677, 0.01444139, 0.02041631, 0.02928188,
        0.01881423, 0.02746681, 0.03195811, 0.01682951, 0.02033133,
        0.022804  , 0.03133063, 0.03071251]))
```